

# CS3331 Concurrent Computing Exam 1 Solutions

## Spring 2018

### 1. Basic Concepts

- (a) [10 points] Explain what are CPU modes. Explain their uses. How does the CPU know what mode it is in? **There are three questions.**

**Answer:** The following has the answers.

- CPU modes are operating modes of the CPU. Modern CPUs have two execution modes: the user mode and the supervisor (or system, kernel, privileged) mode, controlled by a mode bit.
- The OS runs in the supervisor mode and all user programs run in the user mode. Some instructions that may do harm to the OS (e.g., I/O and CPU mode change) are privileged instructions. Privileged instructions, for most cases, can only be used in the supervisor model. When execution switches to the OS (*resp.*, a user program), execution mode is changed to the supervisor (*resp.*, user) mode.
- A mode bit can be set by the operating system, indicating the current CPU mode.

See page 5 of 02-Hardware-OS.pdf. ■

- (b) [10 points] What is an atomic instruction? What would happen if multiple CPUs/cores execute their atomic instructions? **Make sure you will explain atomic instructions fully. Otherwise, you may receive a lower or very low score.**

**Answer:** An atomic instruction is a machine instruction that executes as one *uninterruptible* unit without interleaving and cannot be split by other instructions. When an atomic instruction is recognized by the CPU, we have the following:

- All other instructions being executed in various stages by the CPUs are suspended (and perhaps re-issued later) until this atomic instruction finishes.
- No interrupts can occur.

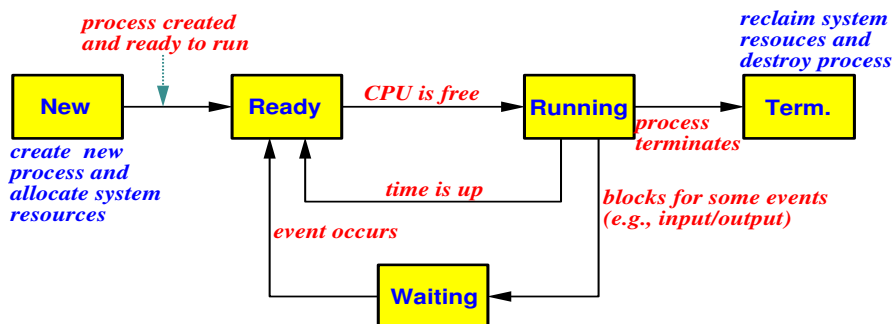
If two such instructions are issued at the same time on different CPUs/cores, they will be executed sequentially in an arbitrary order determined by the hardware.

See pp. 12–13 of 02-Hardware-OS.pdf. ■

### 2. Processes

- (a) [10 points] Draw the state diagram of a process from its creation to termination, including all transitions. Make sure you will elaborate **every state** and **every transition** in the diagram.

**Answer:** The following state diagram is taken from my class note.



There are five states: **new**, **ready**, **running**, **waiting**, and **terminated**.

- **New:** The process is being created.
- **Ready:** The process has everything but the CPU, and is waiting to be assigned to a processor.
- **Running:** The process is executing on a CPU.
- **Waiting:** The process is waiting for some event to occur (*e.g.*, I/O completion or some resource).
- **Terminated:** The process has finished execution.

The transitions between states are as follows:

- **New→Ready:** The process has been created and is ready to run.
- **Ready→Running:** The process is selected by the CPU scheduler and runs on a CPU/core.
- **Running→Ready:** An interrupt has occurred forcing the process to wait for the CPU.
- **Running→Waiting:** The process must wait for an event (*e.g.*, I/O completion or a resource).
- **Waiting→Ready:** The event the process is waiting has occurred, and the process is now ready for execution.
- **Running→Terminated:** The process exits.

See pp. 5–6 of 03-Process.pdf. ■

- (b) [10 points] What is a *context*? Provide a detail description of *all* activities of a *context switch*.

**Answer:** A process needs some system resources (*e.g.*, memory and files) to run properly. These system resources and other information of a process include process ID, process state, registers, memory areas (for instructions, local and global variables, stack and so on), various tables (*e.g.*, PCB), a program counter to indicate the next instruction to be executed, etc. They form the *environment* or *context* of a process. The steps of switching process *A* to process *B* are as follows:

- The operating system suspends *A*'s execution. A CPU mode switch may be needed.
- Transfer the control to the CPU scheduler.
- Save *A*'s context to its PCB and other tables.
- Load *B*'s context to register, etc. from *B*'s PCB.
- Resume *B*'s execution of the instruction at *B*'s program counter. A CPU mode switch may be needed.


See pp. 10–11 of 03-Process.pdf. ■

### 3. Threads

- (a) [10 points] Do your best to define and compare the concepts of *process*, *thread* and *fiber*. In particular, how are they scheduled? **You have to define each first, followed by a comparison of these three. You are expected to explain the way of scheduling each. Without doing so, you will receive no credit for this problem.**

**Answer:** A *process* is the execution of a program; a *thread* (*i.e.*, *lightweight process*), is a unit of CPU execution that is created by a process; and a *fiber* is a lightweight thread, created by a thread, just like a thread being a lightweight process. A comparison can be based on hierarchy, resource usage and scheduling as follows:

- **Hierarchy:** Processes create threads, and threads create fibers. Thus, a process contains the threads it creates, and a thread contains the fibers it creates. This hierarchy is shown below:



- Resource Sharing:** A *process* must acquire all resources for that process to run properly. These resources include files, memory, registers, etc.  
 A *thread* is a lightweight process. All peer threads created within a process share with each other files opened by the containing process, memory allocated to the containing process, etc. Thus, a thread has a program counter, a register set, and a stack, and shares the resources acquired by the containing process with other peer threads.  
 A *fiber* also has a program counter, a subset of the registers, and a stack. A fiber shares all resources acquired by the containing threads with other peer fibers.
- Scheduling:** *processes* are scheduled by the CPU scheduling in the kernel. *Threads* can be user-level threads that are scheduled by a thread scheduler built into a thread library in a user address space, or can be kernel-supported threads that are scheduled by the thread scheduler in the kernel. *Fibers*, on the other hand, are scheduled by a co-operative policy. In other words, a fiber gives up the CPU in a voluntary way, usually through the execution of the `YIELD` statement.



#### 4. Synchronization

- (a) [10 points] Define the meaning of a *race condition*? Answer the question first and use execution sequences with a clear and convincing argument to illustrate your answer. **You must explain step-by-step why your example causes a race condition.**

**Answer:** A *race condition* is a situation in which more than one processes or threads manipulate a shared resource concurrently, and the result depends on the order of execution.

The following is a simple counter updating example discussed in class. The value of `count` may be 9, 10 or 11, depending on the order of execution of the **machine instructions** of `count++` and `count--`.

```

int          count = 10; // shared variable

Process 1          Process 2
count++;           count--;

```

The following execution sequence shows a race condition. Two processes run concurrently (condition 1). Both processes access the shared variable `count` concurrently (condition 2) because `count` is accessed in an interleaved way. Finally, the computation result depends on the order of execution of the `SAVE` instructions (condition 3). The execution sequence below shows the result being 9; however, switching the two `SAVE` instructions yields 11. Since all conditions are met, we have a race condition. **Note that you have to provide TWO execution sequences, one for each possible result, to justify the existence of a race condition.**

Thread_1	Thread_2	Comment
do something	do something	count = 10 initially
LOAD count		Thread_1 executes count++
ADD #1		
	LOAD count	Thread_2 executes count--
	SUB #1	
SAVE count		count is 11 in memory
	SAVE count	Now, count is 9 in memory


Stating that “count++ followed by count--” or “count-- followed by count++”, even using machine instructions, produces different results and hence a race condition is **incomplete**, because the two processes do not access the shared variable `count` concurrently. Note that the use of higher-level language statement interleaved execution may not reveal the key concept of “sharing” as discussed in class. Therefore, use instruction level interleaved instead.

See pp. 5–12 of 05-Sync-Basics.pdf. ■

## 5. Problem Solving:

- (a) [10 points] Design a C program segment so that the `main()` creates two child processes with `fork()`, each of these two child processes creates two child processes, etc. such that the parent-child relationship is a perfectly balanced binary tree of depth  $n$  with `main()` at the root. The depth  $n$  have already been stored a valid positive integer. The `main()` prints its PID, and each child process prints its PID and its parent's PID.

The following diagram shows a tree of processes of depth 3 (i.e.,  $n = 3$ ). Your program segment must be correct for any valid value of  $n > 0$ . Only providing a program segment for  $n = 3$  will receive **zero** point. To save your time, you do not have to perform error checking. However, proper wait and exit are expected.



**Answer:** The parent forks two child processes and waits for them. Each child process prints the needed information, and loops back to fork its own child processes. However, if a child process is a leaf node, it just prints and exits. This idea is illustrated in the code segment below.

```
#include <stdio.h>

#define PRINT { printf("My PID = %ld    My PPID = %ld\n", getpid(), getppid()); }

int main(int argc, char **argv)
{
    int i, n;

    printf("main()'s PID = %ld\n\n", getpid());
    n = atoi(argv[1])-1;
    for (i = 1; i <= n; i++)    // for each level
        if (fork() > 0)        // parent forks the 1st child
            if (fork() > 0) {    // parent forks the 2nd child
                wait(NULL);      // parent waits them to complete
                wait(NULL);
                exit(0);          // and exit
            }
        else {                  // the 2nd child process
            PRINT                // print the needed info
            if (i == n)          // if it is a leaf node
                exit(0);         // exit w/o coming back to fork
        }
    else {                      // the 1st child process
        PRINT                   // print the needed info
        if (i == n)             // if it is a leaf node
            exit(0);             // exit w/o coming back to fork
    }
}
```

You may also do the same using recursion as binary tree creation is basically recursive. Note that a process prints only if the level number is larger than 1. If the level number is 1, printing `getpid()` and `getppid()` will print the PID of the `main()` and the PID of the shell.

```

#include <stdio.h>

#define PRINT    { printf("My PID = %ld    My PPID = %ld\n", getpid(), getppid()); }

void pCreate(int i, int n)
{
    if (i < n) {
        // still have level to go
        if (fork() > 0) {
            // fork the 1st child
            if (fork() > 0) {
                // fork the 2nd child
                if (i > 1)
                    // if this is higher than level 1
                    PRINT
                    // print
                wait(NULL);
                // wait for both child processes
                wait(NULL);
                exit(0);
                // exit
            }
            else
                // the 2nd child goes for level i+1
                pCreate(i+1, n);
        }
        else
            // the 1st child goes for level i+1
            pCreate(i+1, n);
    }
    else {
        // leaf node
        PRINT
        // print info only w/o forking
        exit(0);
    }
}

int main(int argc, char **argv)
{
    int i, n;

    printf("main()'s PID = %ld\n\n", getpid());
    n = atoi(argv[1]);
    pCreate(1, n);
    wait(NULL);
}

```

There are two commonly seen problems in your solutions. **First**, no `wait()` statements were included. If you do not include `wait()` statements, some child processes can become orphans whose parent is the `init` process with PID being 1, and, as a result, the parent-child relationship is incorrect. **Second**, even though you are allowed to use `printf()`, separating the parent PID and the child PID into two `printf()` calls is an incorrect approach. It is because the two `printf()` calls may prints their output lines scattered all over in the output report. As a result, the parent-child relationship is, again, not shown properly. ■

- (b) [15 points] Consider the following two processes, *A* and *B*, to be run concurrently using a shared memory for the `int` variable `x`.

Process A	Process B
-----	-----
for (i = 1; i <= 2; i++)	x = 2*x;
x++;	

Assume that `x` is initialized to 0, and `x` must be loaded into a register before further computations can take place. What are **all possible** values of `x` after both processes have terminated. **You must use clear step-by-step execution sequences of the above processes with a convincing**

**argument. Any vague and unconvincing argument receives no points.**

**Answer:** Obviously, the answer must be in the range of 0 and 4. It is non-negative, because the initial value is 0 and no subtraction is used. It cannot be larger than 4, because the two  $x++$  statements and  $x = 2*x$  together can at most double the value of  $x$  twice.

The easiest answers are 2, 3 and 4 if  $x = 2*x$  executes before, between and after the two  $x++$  statements, respectively. The following shows the possible execution sequences.

$x = 2*x$ is before both $x++$		
Process 1	Process 2	$x$ in memory
	$x = 2*x$	0
$x++$		1
$x++$		2

$x = 2*x$ is between the two $x++$		
Process 1	Process 2	$x$ in memory
$x++$		1
	$x = 2*x$	2
$x++$		3

$x = 2*x$ is after both $x++$		
Process 1	Process 2	$x$ in memory
$x++$		1
$x++$		2
	$x = 2*x$	4

The situation is a bit more complex with instruction interleaving. Process B's  $x = 2*x$  may be translated to the following machine instructions:

```
LOAD x
MUL #2
SAVE x
```

Because the `LOAD` retrieves the value of  $x$ , and the `SAVE` may change the current value of  $x$ , the results depend on the positions of `LOAD` and `SAVE`. The following shows the result being 0. In this case, `LOAD` loads 0 *before* both  $x++$  statements, and the result 0 is saved *after* both  $x++$  statements.

Process 1	Process 2	$x$ in memory	Comments
	<code>LOAD x</code>	0	Load $x = 2$ into register
	<code>MUL #2</code>	0	Process 2's register is 0
$x++$		1	Process 1 adds 1 to $x$
$x++$		2	Process 1 adds 1 to $x$
	<code>SAVE x</code>	0	Process 2 saves 0 to $x$

If the `SAVE` executes between the two  $x++$  statements, the result is 1.

Process 1	Process 2	$x$ in memory	Comments
	<code>LOAD x</code>	0	Load $x = 2$ into register
	<code>MUL #2</code>	0	Process 2's register is 0
$x++$		1	Process 1 adds 1 to $x$
	<code>SAVE x</code>	0	Process 2 saves 0 to $x$
$x++$		1	Process 1 adds 1 to $x$

You may try other instruction interleaving possibilities and the answers should still be in the range of 0 and 4. ■

- (c) [15 points] Consider the following solution to the mutual exclusion problem for two processes  $P_1$  and  $P_2$ . This solution uses two global int variables,  $x$  and  $y$ . Both  $x$  and  $y$  are initialized to 0.

```
int x = 0, y = 0;
```

#### Process 1

```
1 START:
2   x = 1;
3   if (y != 0) {
4       repeat until (y == 0);
5       goto START;
6   }
7   y = 1;
8   if (x != 1) {
9       y = 0;
10      repeat until (x == 0);
11      goto START;
12  }
    // critical section
13  x = y = 0;
```

#### Process 2

```
START:                                // All start from here
x = 2;                                // set my ID to x
if (y != 0) {                          // if y is non-zero
    repeat until (y == 0);             // wait until y = 0
    goto START;                       // then try again
}                                      // second section
y = 1;                                // set y to 1
if (x != 2) {                          // if x is not my ID
    y = 0;                             // set y to 0
    repeat until (x == 0);             // wait until x = 0
    goto START;                       // then try again
}
    // critical section
x = y = 0;                            // set x and y to 0
```

Prove rigorously that this solution satisfies the mutual exclusion condition. *You will receive **zero** point if (1) you prove by example, or (2) your proof is vague and/or unconvincing.*

**Answer:** We shall prove the mutual exclusion property by contradiction. Consider process  $P_1$  first. If  $P_1$  is in its critical section, its execution must have set  $x$  to 1; passed the first if statement which does not modify the value of  $x$ ; set  $y$  to 1; and seen  $x \neq 1$  being false. Therefore, if  $P_1$  is in its critical section,  $x$  and  $y$  must both be 1. By the same reason, if  $P_2$  is in its critical section,  $x$  and  $y$  must be 2 and 1, respectively. Now, if  $P_1$  and  $P_2$  are **both** in their critical sections,  $x$  must be both 1 and 2. This is impossible because a variable can only hold one value. As a result, the assumption that  $P_1$  and  $P_2$  are both in their critical sections cannot hold, and, the mutual exclusion condition is satisfied. ■