

# CS3331 Concurrent Computing Exam 2 Solutions

## Spring 2014

### 1. Synchronization Basics

- (a) [15 points] Consider the following solution to the mutual exclusion problem for two processes  $P_0$  and  $P_1$ . A process can be making a request REQUESTING, executing in the critical section IN\_CS, or having nothing to do with the critical section OUT\_CS. This status information, which is represented by an int, is saved in `flag[i]` of process  $P_i$ . Moreover, variable `turn` is initialized elsewhere to be 0 or 1. Note that `flag[]` and `turn` are global variables shared by both  $P_0$  and  $P_1$ .

```
int    flag[2];    // global flags
int    turn;       // global turn variable, initialized to 0 or 1

Process i (i = 0 or 1)

// Enter Protocol
repeat                                     // repeat the following
    flag[i] = REQUESTING;                 // making a request to enter
    while (turn != i && flag[j] != OUT_CS) // as long as it is not my turn and
        ;                                // the other is not out, wait
    flag[i] = IN_CS;                      // OK, I am in (well, maybe); but,
until flag[j] != IN_CS;                  // must wait until the other is not in
turn = i;                                // the other is out and it is my turn!

// critical section

// Exit Protocol
turn = j;                                // yield the CS to the other
flag[i] = OUT_CS;                        // I am out of the CS
```

Prove rigorously that this solution satisfies the mutual exclusion condition. *You will receive **zero** point if (1) you prove by example, or (2) your proof is vague and/or not convincing.*

**Answer:** A process that enters its critical section must first set `flag[i] = IN_CS` and then see `flag[j] != IN_CS` being true at the end of the `repeat-until` loop. Therefore, if process  $P_0$  is in the critical section, it had executed `flag[0] = IN_CS` followed by seeing `flag[1] != IN_CS`. By the same reason, if process  $P_1$  is in the critical section, it had executed `flag[1] = IN_CS` followed by seeing `flag[0] != IN_CS`. Consequently, if  $P_0$  and  $P_1$  are both in the critical section, we have `flag[0] = IN_CS` and `flag[1] != IN_CS` (from  $P_0$ 's point of view) and `flag[1] = IN_CS` and `flag[0] != IN_CS` (from  $P_1$ 's point of view). As a result, `flag[0]` and `flag[1]` are equal to `IN_CS` and not equal to `IN_CS` at the same time. This is impossible, and the mutual exclusion condition holds.

See page 8 of 06-Sync-Soft-Hardware.pdf. This is exactly the same technique we used to show that Peterson's algorithm satisfies the mutual exclusion condition. ■

- (b) [10 points]\* Define the meaning of a *race condition*? Answer the question first and use an execution sequence with a clear and convincing argument to illustrate your answer. **You must explain step-by-step why your example causes a race condition.**

**Answer:** A *race condition* is a situation in which *more than one* processes or threads access a shared resource *concurrently*, and the result depends on *the order of execution*.

The following is a simple counter updating example discussed in class. The value of `count` may be 9, 10 or 11, depending on the order of execution of the machine instructions of `count++` and `count--`.

```
int    count = 10; // shared variable

Process 1          Process 2

count++;           count--;
```

The following execution sequence shows a race condition. Two processes run concurrently (condition 1). Both processes access the shared variable `count` at the same time (condition 2). Finally, the computation result depends on the order of execution of the `SAVE` instructions (condition 3). The execution sequence below shows the result being 9; however, switching the two `SAVE` instructions yields 11. Since all conditions are met, we have a race condition.

Thread_1	Thread_2	Comment
do something	do something	<code>count = 10</code> initially
LOAD <code>count</code>		Thread_1 executes <code>count++</code>
ADD #1		
	LOAD <code>count</code>	Thread_2 executes <code>count--</code>
	SUB #1	
SAVE <code>count</code>		<code>count</code> is 11 in memory
	SAVE <code>count</code>	Now, <code>count</code> is 9 in memory

Stating that “`count++` followed by `count--`” or “`count--` followed by `count++`” produces different results and hence a race condition is at least **incomplete**, because the two processes do not access the shared variable `count` at the same time (*i.e.*, condition 2).

See page 5 to page 9 of 05-Sync-Basics.pdf. ■

## 2. Semaphores

- (a) [10 points] The semaphore methods `Wait()` and `Signal()` must be atomic to ensure a correct implementation of mutual exclusion. Use execution sequences to show that if `Wait()` is not atomic then mutual exclusion cannot be maintained.

**Answer:** If `Wait()` is not atomic, its execution may be switched in the middle. If this happens, mutual exclusion will not be maintained. Consider the following solution to the critical section problem:

```
Semaphore S = 1;
```

<pre>Process A ----- Wait(S);     // in critical section Signal(S);</pre>	<pre>Process B ----- Wait(S); Signal(S);</pre>
---	--

The following is a possible execution sequence, where `Count = 1` is the internal counter variable of the involved semaphore `S`.

Process A	Process B	Count	Comment
		1	Initial value
LOAD <code>Count</code>		1	A executes <code>Count--</code> of <code>Wait()</code>
SUB #1		1	
	LOAD <code>Count</code>	1	B executes <code>Count--</code> of <code>Wait()</code>
	SUB #1	1	
	SAVE <code>Count</code>	0	B finishes <code>Count--</code>
SAVE <code>Count</code>		0	A finishes <code>Count--</code>
if ( <code>Count &lt; 0</code> )		0	It is false for A
	if ( <code>Count &lt; 0</code> )	0	It is false for B
Both A and B enter the critical section			

Note that this question asks you to demonstrate a violation of mutual exclusion. Consequently, you receive low grade if your demonstration is not a violation of mutual exclusion.

This problem was assigned as an exercise in class. ■

- (b) [10 points] Three ingredients are needed to make a cigarette: tobacco, paper and matches. An agent has an infinite supply of all three. Each of the three smokers has an infinite supply of one ingredient only. That is, one of them has tobacco, the second has paper, and the third has matches. The following solution uses three semaphores, each of which represents an ingredient, and a fourth one to control the table. A smoker waits for the needed ingredients on the corresponding semaphores, signals the table semaphore to tell the agent that the table has been cleared, and smokers for a while.

```
Semaphore Table = 0;           // table semaphore
Semaphore Sem[3] = { 0, 0, 0 }; // ingredient semaphores

int TOBACCO = 0, PAPER = 1, MATCHES = 2

Smoker_Tobacco      Smoker_Paper      Smoker_Matches

while (1) {           while (1) {           while (1) {
    // other work
    Sem[PAPER].Wait();   Sem[TOBACCO].Wait();   Sem[TOBACCO].Wait();
    Sem[MATCHES].Wait(); Sem[MATCHES].Wait();   Sem[PAPER].Wait();
    Table.Signal();      Table.Signal();      Table.Signal();
    // smoke
}                       }                       }
```

The agent adds two randomly selected different ingredients on the table, and signals the corresponding semaphores. This process continues forever.

```
while (1) {
    // generate two different random integers in the range of 0 and 2,
    // say X and Y
    Sem[X].Signal();
    Sem[Y].Signal();
    Table.Wait();
    // do some other tasks
}
```

Show, using execution sequences, that this solution can have a deadlock. **You will receive zero point if you do not use valid execution sequences.**

**Answer:** This is a simple problem and was discussed in class when we talked about the smokers problem. The following shows a possible execution sequence:

Smoker 1	Smoker 2	Smoker 3	Agent	Comment
Wait on PAPER	Wait on TOBACCO	Wait on TOBACCO		All three smokers wait
			Signal TOBACCO	Tobacco available
		Wait in PAPER		Smoker 3 released
			Signal PAPER	Paper available
Wait on MATCHES				Smoker 1 released
			Wait on TABLE	Agent blocks and waits on TABLE
All smokers and the agent blocks				Deadlock occurs

There are other execution sequences based on the same idea. ■

- (c) [15 points] A programmer used two semaphores to design a class **Barrier**, a constructor, and method **Barrier\_wait()** that fulfills the following specification:

- The constructor **Barrier(int n)** takes a positive integer argument **n**, and initializes a private **int** variable in class **Barrier** to have the value of **n**.
- Method **Barrier\_wait(void)** takes no argument. A thread that calls **Barrier\_wait()** blocks if the number of threads being blocked is less than **n-1**, where **n** is the initialization value and will not change in the execution of the program. Then, the **n-th** calling thread releases all **n-1** blocked threads and all **n** threads continue.

This programmer came up with the following solution. However, he found his solution could react strangely because sometimes the same thread may be released multiple times in the same batch. Of course, this is wrong.

```

Semaphore Mutex      = 1;
Semaphore WaitingList = 0;

Barrier::Barrier_wait()
{
    int i;

    Mutex.Wait();           // lock the counter
    if (count == Total-1) { // I am the n-th one
        count = 0;          // reset counter
        Mutex.Signal();     // release the lock
        for (i=1; i<=Total-1; i++) // release all waiting threads
            WaitingList.Signal();
    }                       // I am done
    else {                  // otherwise, I am not the last one
        count++;            // one more waiting threads
        Mutex.Signal();     // release the mutex lock
        WaitingList.Wait(); // block myself
    }
}

```

Help this programmer pinpoint the problem with an execution sequence plus a convincing explanation.

**Answer:** This solution does not have a mechanism to prevent fast-runners from coming back, and, as a result, a thread just released may come back and get released again.

If there are  $n - 1$  threads waiting on semaphore `WaitingList`, where  $n$  is `Total` in the program. Then, thread  $T_n$  comes, locks `Mutex`, sets `count` to 0, unlocks `Mutex`, and starts signaling semaphore `WaitingList`. Right after a thread, say  $T_i$ , is released from semaphore `WaitingList` and before  $T_n$  can signal the second time,  $T_n$  is switched out and  $T_i$  is switched in. As a result,  $T_i$  runs, comes back, sees `count` being 0, increases `count` by 1, and waits on semaphore `WaitingList` again. Now, thread  $T_n$  resumes, and signals the second time. Since there is no particular order for releasing threads from a semaphore,  $T_i$  may be released the second time.

The following is an execution sequence:

Thread $T_i$	Thread $T_n$	count	Comment
		$n - 1$	$n - 1$ threads waiting
	<code>Mutex.Wait()</code>	$n - 1$	$T_n$ arrives
	<code>if ...</code>	$n - 1$	$T_n$ sees a full count
	<code>count = 0</code>	0	$T_n$ resets <code>count</code>
	<code>WaitingList.Signal()</code>	0	$T_n$ signals once
<code>Mutex.Wait()</code>		0	$T_i$ is released and comes back
<code>count++</code>		1	$T_i$ increases <code>count</code>
<code>Mutex.Signal()</code>		1	$T_i$ releases <code>count</code>
<code>WaitingList.Wait()</code>		1	$T_i$ blocks
	<code>WaitingList.Signal()</code>	1	$T_n$ signals the second time
Thread $T_i$ may be released again!			

In this way, the fast runner  $T_i$  is released twice. ■

### 3. Problem Solving:

- (a) [20 points] A multithreaded program has two global arrays and a number of threads that execute concurrently. The following shows the global arrays, where `n` is a constant defined elsewhere (*e.g.*, in a `#define`):

```
int a[n], b[n];
```

Thread  $T_i$  ( $0 < i \leq n-1$ ) runs the following (pseudo-) code, where function  $f()$  takes two integer arguments and returns an integer, and function  $g()$  takes one integer argument and returns an integer. Functions  $f()$  and  $g()$  do not use any global variable.

```
while (not done) {
    a[i] = f(a[i], a[i-1]);
    b[i] = g(a[i]);
}
```

More precisely, thread  $T_i$  passes the value of  $a[i-1]$  computed by  $T_{i-1}$  and the value of  $a[i]$  computed by  $T_i$  to function  $f()$  to compute the new value for  $a[i]$ , which is then passed to function  $g()$  to compute  $b[i]$ .

Declare semaphores with initial values, and add `Wait()` and `Signal()` calls to thread  $T_i$  so that it will compute the result correctly. Your implementation should not have any busy waiting, race condition, and deadlock, and should aim for maximum parallelism.

**A convincing correctness argument is needed. Otherwise, you will receive no credit for this problem.**

**Answer:** If you look at the code carefully, you will see that thread  $T_i$  ( $1 < T_i < n-1$ ) shares  $a[i-1]$  and  $a[i]$  with  $T_{i-1}$  and  $T_{i+1}$ , respectively. Therefore,  $T_i$  must wait until  $T_{i-1}$  completes its update to  $a[i-1]$  before using it in the computation of  $a[i] = f(a[i], a[i-1])$ . Similarly,  $T_i$  must wait until  $T_{i+1}$  finishes using  $a[i]$  before computing a new value for  $a[i]$ . To this end, we need a semaphore  $s[i-1]$  for protecting  $a[i-1]$  and a semaphore  $s[i]$  for protecting  $a[i]$ . Isn't this very similar to the philosophers problem if you considered  $T_i$  as a philosopher and  $a[i-1]$  and  $a[i]$  as  $T_i$ 's chopsticks? The major difference is that the philosophers are "circular" while the  $T_i$ 's are "linear."

Next, we consider  $T_1$  and  $T_{n-1}$ .  $T_1$  uses  $a[0]$  and  $a[1]$ . Since no thread updates  $a[0]$ ,  $T_1$  only needs a semaphore  $s[1]$  to protect  $a[1]$ . Similarly, thread  $T_{n-1}$  uses  $a[n-1]$  and  $a[n-2]$ . Since there is no thread using  $a[n-1]$  other than  $T_{n-1}$  itself,  $T_{n-1}$  only needs a semaphore  $s[n-2]$  to protect  $a[n-2]$ .

The following is a possible solution. Note that  $b[i]$  is not protected by any semaphore because (1) only  $T_i$  uses  $b[i]$  and (2) only  $T_i$  writes into  $a[i]$ , and once  $a[i]$  is correctly computed one can compute  $b[i]$  correctly.

```
Sem s[1..n-1] = { 1, 1, 1, ..., 1 }; // all semaphores = 1 for mutual exclusion
```

<pre>Thread-1: while (not done) {     s[1].Wait();     a[1] = f(a[0], a[1]);     s[1].Signal();     b[1] = g(a[1]); }</pre>	<pre>Thread-i: while (not done) {     s[i-1].Wait();     s[i].Wait();     a[i] = f(a[i-1], a[i]);     s[i].Signal();     s[i-1].Signal();     b[i] = g(a[i]); }</pre>	<pre>Thread-(n-1): while (not done) {     s[n-2].Wait();     a[n-1] = f(a[n-2], a[n-1]);     s[n-2].Signal();     b[n-1] = g(a[n-1]); }</pre>
---	---	---

Obviously, there is no busy waiting. The above solution has no race condition either because each shared data item (*i.e.*,  $a[i-1]$  and  $a[i]$  for thread  $T_i$ ) is protected by mutual exclusion.

This solution is deadlock free. What we have to show is that no semaphore would block threads indefinitely. Consider thread  $T_i$ , which may be blocked on semaphores  $s[i-1]$  or  $s[i]$ . If  $T_i$  is blocked on semaphore  $s[i-1]$ , this means thread  $T_{i-1}$  is now in its critical section. Thread  $T_{i-1}$  will eventually exit its critical section and execute  $s[i-1].Signal()$ . At this moment, thread  $T_i$  could continue although  $T_{i-1}$  may come back fast enough and execute  $s[i-1].Wait()$  successfully again. Whatever the case is,  $T_i$  has the chance to continue, and, the worst case is starvation rather than a deadlock.

If thread  $T_i$  is blocked on semaphore  $s[i]$ , this means thread  $T_{i+1}$  executed its  $s[i].Wait()$ , followed by  $s[i+1].Wait()$ . See the figure above. Thread  $T_{i+1}$  may or may not be blocked by  $s[i+1].Wait()$ . If thread  $T_{i+1}$  is not blocked by  $s[i+1].Wait()$ , it will eventually signal it, allowing  $T_i$  to continue. If thread  $T_{i+1}$  is blocked by  $s[i+1].Wait()$ , it is because thread  $T_{i+2}$



- (b) [20 points] A unisex bathroom is shared by men and women. A man or a woman may be using the room, waiting to use the room, or doing something else. They work, use the bathroom and come back to work. The rule of using the bathroom is very simple: *there must never be a man and a woman in the room at the same time; however, people with the same gender can use the room at the same time.*

**Man Thread**

```
void Man(void)
{
    while (1) {
        // working
        // use the bathroom
    }
}
```

**Woman Thread**

```
void Woman(void)
{
    while (1) {
        // working
        // use the bathroom
    }
}
```

Declare semaphores and other variables with initial values, and add `Wait()` and `Signal()` calls to the threads so that the man threads and woman threads will run properly and meet the requirement. Your implementation should not have any busy waiting, race condition, and deadlock, and should aim for maximum parallelism.

**A convincing correctness argument is needed. Otherwise, you will receive no credit for this problem.**

**Answer:** This is a simple variation of the reader-priority readers-writers problem. More precisely, we allow the “writers” to write simultaneously. Therefore, the writers have the same structures as the readers. We need to maintain two counters, one for the males `MaleCounter` and the other for the females `FemaleCounter`. Of course, we need two Mutexes `MaleMutex` and `FemaleMutex` for mutual exclusion. In addition, there is a semaphore `BathRoom` to block the males (*resp.*, females) if the room is being used by the females (*reap.*, males). Note that the male thread and female thread are symmetric.

```
int      MaleCounter = 0, FemaleCounter = 0;    // male and female counters
Semaphore MaleMutex = 1, FemaleMutex = 1;      // male and female counters
Semaphore BathRoom = 1;                        // the bathroom is empty initially

Male Thread                                Female Thread

while (1) {                                while(1) {
    // working                               // working

    MaleMutex.Wait();                        FemaleMutex().Wait();    // update counter
    MaleCounter++;                            FemaleCounter--;
    if (MaleCounter == 1)                    if (FemaleCounter == 1) // if I am the first
        BathRoom.Wait();                      BathRoom.Wait();      //   yield to other
    MaleMutex.Signal();                      FemaleMutex.Signal();

    // use the bathroom                       // use the bathroom

    MaleMutex.Wait();                        FemaleMutex.Wait();      // update counter
    MaleCounter--;                            FemaleCounter--;
    if (MaleCounter == 0)                    if (FemaleCounter == 0) // if I am the last one
        BathRoom.Signal();                      BathRoom.Signal();    //   let the other group know
    MaleMutex.Signal();                      FemaleMutex.Signal();

}                                              }
```

Refer to the class note for the solution to the reader-priority version of the readers-writers problem for the details. ■